

1 Introduction

2 Models with Built-In Feature Selection

Many models included in `caret` have built-in feature selection, including `rpart`, `gbm`, `ada`, `glmboost`, `gamboost`, `blackboost`, `ctree`, `sparseLDA`, `sddaLDA`, `sddaQDA`, `glmnet`, `lasso`, `lars`, `spls`, `earth`, `fda`, `bagEarth`, `bagFDA`, `pam` and others. Many of the functions have an ancillary method called `predictors` that returns a vector indicating which predictors were used in the final model.

In many cases, using these models with built-in feature selection will be more efficient than algorithms where the search routine for the right predictors is “outside” the model. Built-in feature selection couples the predictor search algorithm with the parameter fitting and are usually optimize with a single objective function (e.g. error rates or likelihood).

3 Feature Selection Using Search Algorithms

3.1 Searching the Feature Space

Search routines; validation

3.2 Resampling and External Validation

3.3 Backwards Selection

The recursive feature elimination (RFE), a.k.a. backwards selection, can be used to estimate the appropriate number of predictors. First, the algorithm fits the model to all predictors. Each predictor is ranked on how important it is to the model. Let S_i be a sequence of ordered numbers which are candidate values for the number of predictors to retain. At each iteration of feature selection, the S_i top ranked predictors are retained, the model is refit and performance is assessed. The value of S_i with the best performance is chosen and the top S_i predictors are used to fit the final model. Algorithm 1 has a more complete definition.

The algorithm has an optional step where the predictor rankings are recomputed on the model on the reduced feature set. RFE used RFE with random forest and reported that there was XXX.

There are some cases where re-ranking the predictors could help. XXX put this below XXX.

```
Tune/train the model on the training set using all predictors
Calculate model performance
Calculate variable importance or rankings
for Each subset size  $S_i$ ,  $i = 1 \dots S$  do
  Keep the  $S_i$  most important variables
  Tune/train the model on the training set using  $S_i$  predictors
  Calculate model performance
  [Optional] Recalculate the rankings for each predictor
end
Calculate the performance profile over the  $S_i$ 
Determine the appropriate number of predictors and the final ranks of each predictor
Fit the final model based on the optimal  $S_i$ 
```

Algorithm 1: Recursive feature elimination

As previously mentioned, computing the appropriate performance measure can be difficult. To get performance estimates that incorporate the variation due to feature selection, it is suggested that the steps in be “wrapped” inside a layer of resampling (e.g. 10-fold cross-validation). Algorithm 2 shows a version of the algorithm that uses resampling.

While this will provide better estimates of performance, it is more computationally burdensome. For users with access to machines with multiple processors, the first **For** loop in Algorithm 2 can be easily parallelized. Another complication to using resampling is that multiple lists of the “best” predictors are generated at each iteration. At first this may seem like a disadvantage, but it does provide a more probabilistic assessment of predictor importance than a ranking based on a single, fixed data set. At the end of the algorithm, a consensus ranking can be used to determine the best

predictors to retain.

```
for Each Resampling Iteration do
  Partition data into training and test/hold-back set via resampling
  Tune/train the model on the training set using all predictors
  Predict the held-back samples
  Calculate variable importance or rankings
  for Each subset size  $S_i$ ,  $i = 1 \dots S$  do
    Keep the  $S_i$  most important variables
    Tune/train the model on the training set using  $S_i$  predictors
    Predict the held-back samples
    [Optional] Recalculate the rankings for each predictor
  end
end
Calculate the performance profile over the  $S_i$  using the held-back samples
Determine the appropriate number of predictors and the final ranks of each predictor
Fit the final model based on the optimal  $S_i$  using the original training set
```

Algorithm 2: Recursive feature elimination incorporating resampling

4 Recursive Feature Elimination via `caret`

4.1 An Example

To test the algorithm, the “Friedman 1” benchmark problem (Friedman, 1991) was used. there are three informative variables generated with

$$y = 10 \sin(\pi x_1 x_2) + 20(x_3 - 0.5)^2 + 10x_4 + 5x_5 + N(0, \sigma^2)$$

In the simulation that follows:

```
> n <- 100
> p <- 40
> sigma <- 1
> set.seed(1)
> sim <- mlbench.friedman1(n, sd = sigma)
> x <- cbind(sim$x, matrix(rnorm(n * p), nrow = n))
> y <- sim$y
> colnames(x) <- paste("var", 1:ncol(x), sep = "")
```

Of the 50 predictors, there are 45 pure noise variables: 5 are uniform on $[0, 1]$ and 40 are random univariate standard normals.

The predictors are centered and scaled:

```
> normalization <- preProcess(x)
> x <- predict(normalization, x)
> x <- as.data.frame(x)
> subsets <- c(1:5, 10, 15, 20, 25)
```

The simulation will fit models with subset sizes of 25, 20, 15, 10, 5, 4, 3, 2, 1. A linear model and random forests will be used.

```
> set.seed(10)
> ctrl <- rfeControl(functions = lmFuncs, method = "cv", verbose = FALSE,
+   returnResamp = "final")
> lmProfile <- rfe(x, y, sizes = subsets, rfeControl = ctrl)
> print(lmProfile)
```

Recursive feature selection

Outer resampling method was 10 iterations of cross-validation.

Resampling performance over subset size:

Variables	RMSE	Rsquared	RMSESD	RsquaredSD	Selected
1	3.473	0.5285	0.4706	0.1219	
2	3.134	0.6161	0.5937	0.1757	
3	2.954	0.6770	0.9152	0.2242	*
4	3.055	0.6520	0.9889	0.2359	
5	3.229	0.6188	0.8714	0.1966	
10	3.493	0.5549	0.9811	0.2098	
15	3.754	0.5010	1.1806	0.2243	
20	3.893	0.4725	1.0039	0.2026	
25	4.306	0.4009	0.9284	0.1870	
50	4.306	0.4009	0.9284	0.1870	

The top 3 variables (out of 3):

var4, var5, var2

4.2 Helper Functions

4.2.1 the fit function

This function builds the model based on the current data set. The possible arguments are:

- **x**: the current training set of predictor data with the appropriate subset of variables
- **y**: the current outcome data (either a numeric or factor vector)
- **first**: a single logical values for whether the current predictor set has all possible variables
- **...**: optional arguments to pass to the fit function in the call to **rfe**

The **first** argument can be useful. For example, if a random forest model is fit, you may only want the initial model with all predictor variables to be run with **importance = TRUE**.

This function should return a model fit function that can be used for prediction

```
> rffuncs$fit
```

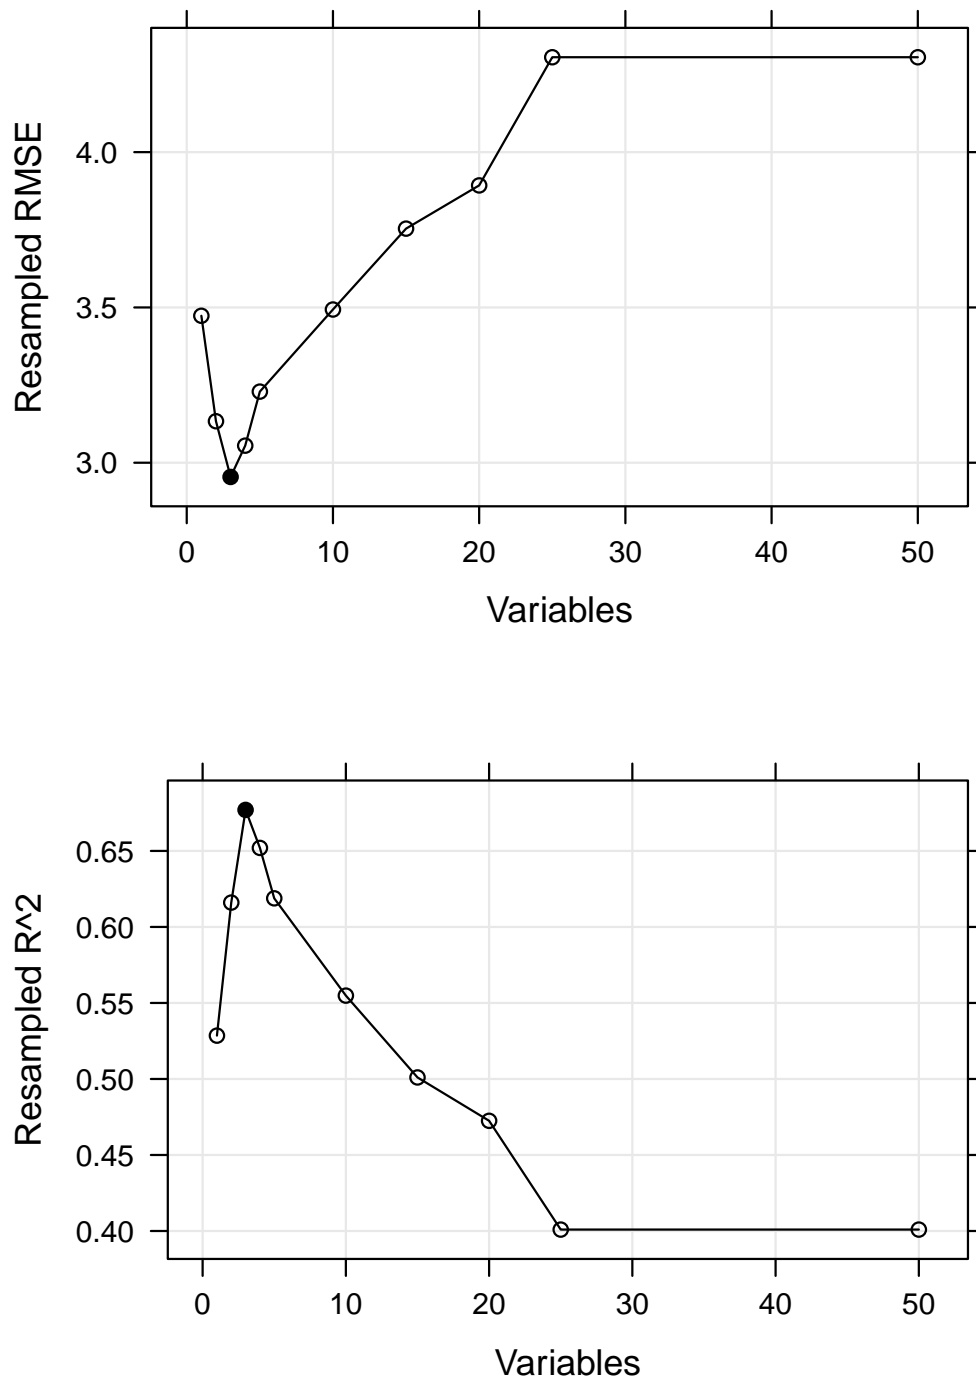


Figure 1: Hold-out performance distributions for four models with built-in feature selection.

```
function (x, y, first, last, ...)  
{  
  library(randomForest)  
  randomForest(x, y, importance = first, ...)  
}  
<environment: namespace:caret>
```

4.2.2 the `pred` function

This function returns a vector of predictions (numeric or factors) from the model. The input arguments must be

- **object**: the model generated by the `fit` function
- **x**: the current set of predictor set for the held-back samples

```
> rfFuncs$pred
```

```
function (object, x)  
{  
  predict(object, x)  
}  
<environment: namespace:caret>
```

4.2.3 the `rank` function

This function should return XXXX.

Inputs are:

- **object**: the model generated by the `fit` function
- **x**: the current set of predictor set for the training samples
- **y**: the current training outcomes

```
> rfFuncs$rank
```

```
function (object, x, y)  
{  
  vimp <- varImp(object)
```

```
if (is.factor(y)) {
  if (all(levels(y) %in% colnames(vimp))) {
    avImp <- apply(vimp[, levels(y), drop = TRUE], 1,
                  mean)
    vimp$Overall <- avImp
  }
}
vimp <- vimp[order(vimp$Overall, decreasing = TRUE), , drop = FALSE]
vimp$var <- rownames(vimp)
vimp
}
<environment: namespace:caret>
```

4.2.4 the selectVar function

Inputs for the function are:

- **y**: a list of variables importance for each resampling iteration and each subset size (generated by the user-defined **rank** function)
- **size**: the subset sized passed into the call to **rfe**

This function should return character string of predictor names (of length **size**) in the order of most important to least important

```
> rfFuncs$selectVar
```

```
function (y, size)
{
  sizes <- unlist(lapply(y[[1]], nrow))
  sizeIndex <- which(size == sizes)
  allImp <- do.call("rbind", lapply(y, function(u, pos) u[[pos]],
    pos = sizeIndex))
  meanImp <- aggregate(allImp[, grep("Overall$", names(allImp))[1]],
    list(var = allImp$var), mean)
  meanImp$imp <- meanImp$x
  counts <- aggregate(allImp[, grep("Overall$", names(allImp))[1]],
    list(var = allImp$var), length)
  counts$pct <- counts$x/length(y)
  counts$x <- NULL
  varInfo <- merge(counts, meanImp)
```



```
varInfo <- varInfo[order(varInfo$pct, varInfo$imp, decreasing = TRUE),
  ]
as.character(varInfo$var[1:size])
}
<environment: namespace:caret>
```

4.2.5 the selectSize function

Inputs for the function are:

- **x**: a matrix with columns for the performance metrics and the number of variables, called **Variables**
- **metric**: a character string of the performance measure to optimize (e.g. RMSE, Accuracy)
- **maximize**: a single logical for whether the metric should be maximized

This function should return an integer that indicates the row of **x** that is optimal

caret comes with two examples runtimes for this purpose: **selectBest**

For example, suppose we have computed the RMSE over a series of variables sizes:

```
> set.seed(10)
> ctrl$functions <- rfFuncs
> rfProfile <- rfe(x, y, sizes = subsets, rfeControl = ctrl)
> print(rfProfile)
```

Recursive feature selection

Outer resampling method was 10 iterations of cross-validation.

Resampling performance over subset size:

Variables	RMSE	Rsquared	RMSESD	RsquaredSD	Selected
1	3.607	0.4670	0.2765	0.16005	
2	3.186	0.6079	0.5151	0.14583	
3	2.779	0.7409	0.3943	0.06699	*
4	2.885	0.7356	0.2721	0.10742	
5	3.177	0.6806	0.4035	0.10557	
10	3.234	0.6726	0.3771	0.11912	
15	3.350	0.6648	0.3780	0.12272	




Figure 2: ads

20	3.415	0.6294	0.3848	0.14043
25	3.588	0.6166	0.3591	0.13230
50	3.565	0.6293	0.3716	0.14347

The top 3 variables (out of 3):

var4, var5, var2

5 Session Information

- R version 2.9.0 Under development (unstable) (2009-01-22 r47686), i386-apple-darwin9.6.0
- Locale: en_US.UTF-8/en_US.UTF-8/C/C/en_US.UTF-8/en_US.UTF-8
- Base packages: base, datasets, graphics, grDevices, grid, methods, splines, stats, tools, utils
- Other packages: caret 4.10, class 7.2-45, e1071 1.5-19, ellipse 0.3-5, gbm 1.6-3, Hmisc 3.5-0, ipred 0.8-6, kernlab 0.9-8, klaR 0.5-8, lattice 0.17-20, MASS 7.2-45, mlbench 1.1-5, nnet 7.2-45, pls 2.1-0, proxy 0.4-1, randomForest 4.5-28, rpart 3.1-42, survival 2.34-1
- Loaded via a namespace (and not attached): cluster 1.11.12

6 References

- Chun, H. and Keles, S. (2007) “Sparse partial least squares for simultaneous dimension reduction and variable selection”, http://www.stat.wisc.edu/~keles/Papers/SPLS_Nov07.pdf.
- Friedman, J. H. (1991) “Multivariate Adaptive Regression Splines (with discussion),” *Annals of Statistics*, 19, 1–141
- Friedman, J. H. (2001) “Greedy Function Approximation: A Gradient Boosting Machine,” *Annals of Statistics*, 29, 1189–1232
- Zou, H. and Hastie, T. (2005) “Regularization and Variable Selection via the Elastic Net,” *Journal of the Royal Statistical Society, Series B*, 67, 301–320.




Figure 3: Hold-out performance distributions for four models with built-in feature selection.